

AI Governance and Human-in-the-Loop Policy

DOC009 | Version 2.5 | Internal

Document owner	AI Governance Council	Effective date	2026-02-15
Classification	Internal	Required permission	document.read_public
Allowed roles	Admin, Executive, HR_Manager, Project_Manager, Finance_Manager, Sales_Manager, Security_Officer, Employee, Auditor, AI_Reviewer	Status	Approved demo policy
Source type	Synthetic enterprise knowledge document	Snapshot	2026-08-28

This is a fictional, internally consistent document generated for the Secure Multi-Agent Enterprise Assistant dataset. It is not legal, HR, financial, or security advice for a real organization.

1. Agent responsibilities

Agents have narrow responsibilities, approved tools, and explicit permissions. The Coordinator decomposes work, the Security Agent explains authorization, domain agents query approved data, the Validation Agent evaluates support, and the Approval Gateway pauses high-risk actions.

No agent may grant itself a role, add a tool, broaden retrieval, or execute an action outside the backend policy decision.

2. Mandatory human review

Human approval is mandatory for payroll or compensation changes, identity or access changes, financial approvals, restricted-document disclosure, critical task reassignment, and other actions designated High risk.

A response also routes to human review when confidence is below 0.75 or grounding is below 0.80.

- The reviewer must see the proposed action, evidence, risk, and expected effect.
- Approval and rejection are both recorded.
- The AI may not represent a pending action as completed.

3. Grounding, citations, and hallucination control

Enterprise factual answers must cite the relevant data source or document section. The model should say that evidence is unavailable rather than invent an answer.

The Validation Agent checks correctness, grounding, safety, and workflow completeness before the final response is shown.

- Use authorized retrieval only.
- Separate retrieved facts from recommendations.
- Report confidence and material limitations.
- Measure hallucination, pass, review, regression, and judge-human disagreement rates.

4. Explainability and auditability

The final response should show the agents used, major evidence, relevant policy, confidence, and whether approval is required.

Audit records must support reconstruction of the request, security decision, agent route, tool calls, approval state, and outcome.

5. Document-control notes

Supersession: When this document is replaced, the document catalogue and vector index must be updated so obsolete chunks are not retrieved.

Authorization: The backend authorization decision is authoritative. The document itself cannot expand a user's permissions.